



## Guaranteed Generation from Large Language Models

Minbeom Kim

Thibaut Thonet

Jos Rozen

Hwaran Lee

Kyomin Jung

Marc Dymetman

### Goals of Guaranteed Generation

1. **100%** strict **guarantees** on constraints satisfaction.
2. **Preserve original distribution** as much as possible.
3. Achieve (1) & (2) with **limited inference costs**.

### Basic Notations

Base LLM:  $a(y)$ , sample  $y \sim a(y)$

The guarantee: binary hard constraint  $b(y) \in \{0, 1\}$

E.g.  $b(y)$  is a toxicity detector (0 means toxic, 1 means non-toxic)

E.g.  $b(y)$  is a verifier of certain keywords

Guaranteed sampler: sampler that **never** violates  $b(y)$

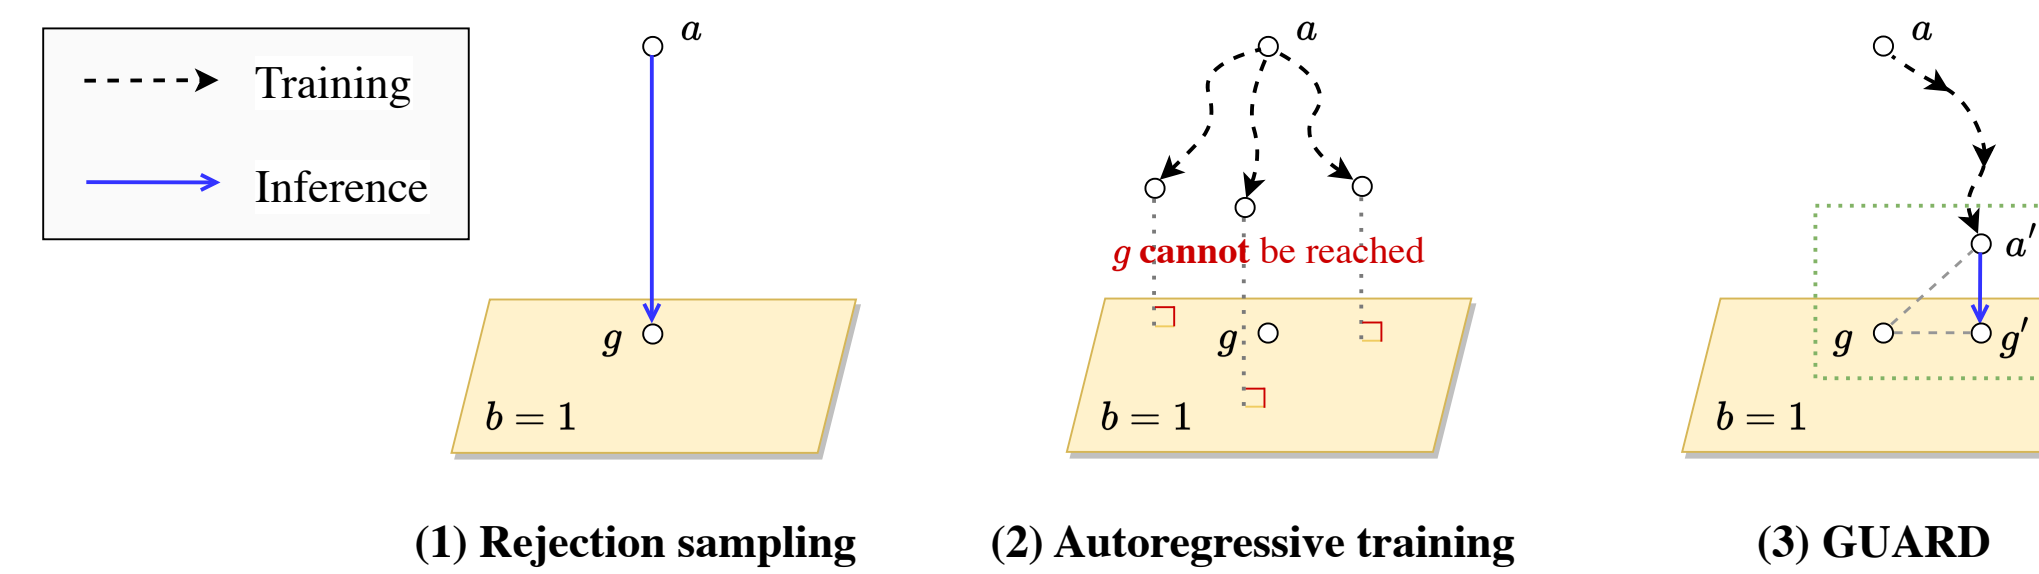
### Formalization of Guaranteed Generation

We define Ideal distribution  $g$  as

$$g(y) = \begin{cases} 0, & b(y) = 0 \\ \frac{1}{Z_{a,b}} a(y), & b(y) = 1 \end{cases}$$

- $g$  is the **guaranteed** distribution,  $a$  conditioned by the fact that  $y$  satisfies  $b$ , i.e.  $g(y) = a(y | [b(y) = 1])$ .
- $g$  is the guaranteed distribution  $p$  that minimizes  $KL(p || a)$ ,  $g$  **preserves the original distribution**  $a$  as much as possible.

### Limitation of Autoregressive Models



**Theorem 1.** However, it is impossible to exactly fit  $g(y)$  with an autoregressive trained model  $a'(y)$ .

**Claim 1.** To achieve this objective, we need **inference-time** Monte-Carlo (MC) methods that exploits the approximation  $a'$  as a *proposal* sampler but only *retains* samples satisfying the constraint  $b$ .

**Claim 2.** When we apply *rejection sampling* to  $a$ , we obtain  $g$  with significant Inference cost. Even with a efficient proposal  $a'$ , a significant inference cost must still be incurred to improve the *quality* of samples via an MCMC.

To address this, we introduce **GUARD** to **amortize** this inference procedure.

### GUARD Framework and its properties

#### Algorithm 1 GUARD sampler

```

1: while True do
2:    $y \sim a'$ 
3:   if  $b(y) = 1$  return  $y$ 

```

**Theorem 2.**  $KL(g || a') = \underbrace{-\log AR_{a'}}_{\text{objective}} + \underbrace{KL(g || g')}_{\text{efficiency}} + \underbrace{KL(g' || a')}_{\text{closeness}}$

? How can we optimize  $KL(g || a')$ ?


**CAP:** prompting  $a$  to be aware of constraint  $b$

**SFT:** sample a lot of  $y \sim a$  and retain only  $b(y) = 1$ . Then fine-tune those on  $a$   
 $\nabla_{\theta} KL(g || a'_{\theta}) = -\mathbb{E}_{y \sim a'_{\theta}} \frac{g(y)}{a'_{\theta}(y)} \nabla_{\theta} \log a'_{\theta}(y)$ , which is *distributional* policy gradient loss,

**DPG:** sample  $y$  from **adaptive** proposal  $a'_{\theta}$  and update  $a'$  with  $\nabla_{\theta} KL(g || a'_{\theta})$ .

Both SFT and DPG is very slow in the early stage... Can we boost up?


**Warm-start DPG:** Fine-tune  $a'$  with  $y \sim a(\cdot | \text{CAP})$  for **skipping early stage**.



### Experiment & Discussion

**Constraint 1:** Lexical constraints “amazing”  $AR_a = 0.23\%$

**Constraint 2:** Story completion from negative opening with “positive” ending  $AR_a = 0.5\%$





- DPG's **adaptive** proposal excels other baselines.
- **Warm-start** helps to skip the inefficient early stage.

• After **GUARD** training,

$AR_{a'}$  of constraint 1: 0.23 %  $\rightarrow$  0.416 % **(180x boost-up!)**

$AR_{a'}$  of constraint 2: 0.5 %  $\rightarrow$  0.306 % **(60x boost-up!)**

**while preserving a high proximity to  $g$**



**Conclusion.** We formalize how to **guarantee** that LLMs perfectly meet specified requirements **without compromising their usefulness**.

Various analysis show minimal distortion from **GUARD** training



GUARD