

Minbeom Kim

✉ research@minbeom.kim | 🏠 https://minbeomkim.github.io/ | 📄 Google Scholar | 🌐 mbkimm

Education

Seoul National University

Ph.D. in Artificial Intelligence

Thesis: Controllable Post-Training of Language Models via Distributional Matching

Best Ph.D Dissertation Award

Advisor: [Kyomin Jung](#)

Seoul, South Korea

Sep. 2021 - Aug. 2026

KAIST

B.E. in Industrial and System Engineering

Daejeon, South Korea

Mar. 2014 - Feb. 2018

Professional Experiences

SK hynix Research

Staff Researcher

Recursive Self-Improvement of AI Agents via Self-evolving Memory

Bundang, South Korea

Jul. 2026 -

Google

Student Researcher (worked on reliable multi-agent AI systems)

Host: [Long T. Le](#)

California, United States

Oct. 2025 - Jun. 2026

LawZero

Research Intern (worked on Scientist AI: probabilistic safety guarantees for superintelligent AI)

Host: [Yoshua Bengio](#)

Montreal, Canada

Jun. 2025 - Sep. 2025

NAVER Labs Europe

Research Intern (worked on theoretical guarantees for constraint satisfaction in generative models)

Host: [Marc Dymetman](#)

Grenoble, France

Apr. 2024 - Sep. 2024

Agency for Defense Development

Research Officer / 1st lieutenant (duration of military service)

I worked on R&D Projects for National Defense Technology. This job also fulfilled my 3-year mandatory military service in South Korea

Daejeon, South.Korea

June. 2018 - May. 2021

Publications

- [1] **Minbeom Kim**, Lesly Miculicich, Bhavana Dalvi Mishra, Mihir Parmar, Phillip Wallis, Bharath Chandrasekhar, Kyomin Jung, Tomas Pfister, Long T. Le, "LiSA: Lifelong Safety Adaptation via Conservative Policy Induction", In: **NeurIPS 2026**
- [2] Jeonghye Kim, Xufang Luo, **Minbeom Kim**, Sangmook Lee, Dongsheng Li, Yuqing Yang, "Understanding Reasoning as Strategic Information Allocation under Uncertainty", **arXiv preprint**
- [3] Sangmook Lee, **Minbeom Kim**, Jeonghye Kim, Dohyung Kim, Sojeong Lee, Kyomin Jung, "Are We Measuring Strategy or Phrasing? The Gap Between Surface- and Approach-Level Diversity in LLM Math Reasoning", In: **EMNLP 2026**
In: **ICML 2026 AI4Math Workshop (Spotlight, Top 2.9%)**
- [4] Jeonghye Kim, Xufang Luo, **Minbeom Kim**, Sangmook Lee, Dohyung Kim, Jiwon Jeon, Dongsheng Li, Yuqing Yang, "Why Does Self-Distillation (Sometimes) Degrade the Reasoning Capability of LLMs?", In: **COLM 2026**

- [5] Kyungmin Min, **Minbeom Kim**, Kang-il Lee, Seunghyun Yoon, Kyomin Jung, "ReflectCap: Detailed Image Captioning with Reflective Multi Agents",
In: **ECCV 2026**
- [6] **Minbeom Kim**, Mihir Parmar, Phillip Wallis, Lesly Miculicich, Kyomin Jung, Krishnamurthy Dj Dvijotham, Long T. Le, Tomas Pfister. "CausalArmor: Efficient Indirect Prompt Injection Guardrails via Causal Attribution",
In: **ICML 2026**
- [7] Dohyung Kim, **Minbeom Kim**, Jeonghye Kim, Sangmook Lee, Sojeong Lee, Kyomin Jung, "Beyond Normalization: Rethinking the Partition Function as a Difficulty Scheduler for RLVR",
In: **ICML 2026**
- [8] Kang-il Lee, Jahyun Koo, Seunghyun Yoon, **Minbeom Kim**, Hyukhun Koh, Dongryeol Lee, Kyomin Jung. "Transductive Program Synthesis",
In: **NeurIPS 2025**
- [9] **Minbeom Kim**, Kang-il Lee, Seongho Joo, Hwaran Lee, Thibaut Thonet, Kyomin Jung. "Drift: Decoding-time Personalized Alignments with *Implicit* User Preferences",
In: **EMNLP 2025**
- [10] Jeonghye Kim*, Sojeong Rhee*, **Minbeom Kim**, Sangmook Lee, Dohyung Lee, Youngchul Sung, Kyomin Jung. "ReflAct: World-Grounded Decision Making in LLM Agents via Goal-State Reflection",
In: **EMNLP 2025**
- [11] Kyeongman Park, **Minbeom Kim**, and Kyomin Jung. "A Character-Centric Creative Story Generation via Imagination",
In: **ACL 2025 Findings**
- [12] **Minbeom Kim**, Thibaut Thonet, Jos Rozen, Hwaran Lee, Kyomin Jung, and Marc Dymetman. "Guaranteed Generation from Large Language Models",
In: **ICLR 2025**
- [13] **Minbeom Kim**, Hwanhee Lee, Joonsuk Park, Hwaran Lee and Kyomin Jung. "AdvisorQA: A Benchmark for Advice-seeking Question Answering with Collective Intelligence",
In: **NAACL 2025**
- [14] Kang-il Lee, **Minbeom Kim**, Seunghyun Yoon, Minsung Kim, Dongryeol Lee, Hyukhun Koh, and Kyomin Jung, "VLind-Bench: Measuring Language Priors in Large Vision-Language Models",
In: **NAACL 2025 Findings**
- [15] Kyungmin Min, **Minbeom Kim**, Kang-il Lee, Dongryeol Lee, and Kyomin Jung, "Mitigating Hallucinations in LVLMs via Summary-Guided Decoding",
In: **NAACL 2025 Findings**
- [16] **Minbeom Kim**, Jahyun Koo, Hwanhee Lee, Joonsuk Park, Hwaran Lee and Kyomin Jung. "LifeTox: Unveiling Implicit Toxicity in Advice-seeking Question Answering",
In: **NAACL 2024 (Oral, Top 5.9%)**
- [17] **Minbeom Kim**, Hwanhee Lee, Kang Min Yoo, Joonsuk Park, Hwaran Lee and Kyomin Jung. "Critic-Guided Decoding for Controlled Text Generation",
In: **ACL 2023 (Findings Spotlights)**
- [18] **Minbeom Kim**, Kyeongha Rho, Yong-duk Kim, Kyomin Jung. "Action-driven Contrastive Representation for Reinforcement Learning",
In: **PLOS ONE 2022**

Awards and Honors

2026	Awards: Best Ph.D Dissertation Award (Seoul National University)	<i>Korea</i>
2024	Awards: Best Summer Intern in Naver Labs Europe	<i>France</i>
2021-2023	Scholarship: Artificial Intelligence Graduate School Program (Seoul National University)	<i>Korea</i>
2022	Awards: Creative AI Research Competition, 2rd (Seoul National University)	<i>Korea</i>

Academic Activities

Academic Services	Area Chair at ACL ARR 2026 ~
	ACL ARR, ICLR, ICML, NeurIPS 2026
	ACL ARR, ICLR, ICML, NeurIPS 2025
	ACL ARR, NeurIPS 2024
	Several Times Nominated as a Great Reviewer at ACL ARR
	Safety Guaranteed Generation from Large Language Models, 2025, Mila Quebec, Canada
Teaching Assistant	AI Retreat, 2023, Seoul National University, South Korea
	Programming Methodology (Spring 2022, 2025), Seoul National University
	IoT, Artificial intelligence, and Big data (IAB) (Fall 2022, 2023), Seoul National University